

How to do unsupervised learning with an autoencoder

By WJG. Took the code from <https://blog.keras.io/building-autoencoders-in-keras.html>
(<https://blog.keras.io/building-autoencoders-in-keras.html>)

1. Create model

- The autoencoder architecture is the following: Input --> 32FC layer --> Output FC means 'fully connected layer'.
- In the code, the layers are coded as follows:
 - First part (Input --> 32FC layer) is denoted as `encoder`;
 - Second part (32D layer --> Output) is denoted as `decoder`.
 - Concatenating the encoder and decoder yields the autoencoder model.
- The 32FC layer (i.e., output of the encoder) is treated as the encoded representation (aka feature vector).
- The decoder can be seen as a generative model taking as input 32D vectors.

In [1]:

```
from keras.layers import Input, Dense
from keras.models import Model

# this is the size of our encoded representations
encoding_dim = 32  # 32 floats -> compression of factor 24.5, assuming the input is 784 floats

# this is our input placeholder
input_img = Input(shape=(784,))
# "encoded" is the encoded representation of the input
encoded = Dense(encoding_dim, activation='relu')(input_img)
# "decoded" is the lossy reconstruction of the input
decoded = Dense(784, activation='sigmoid')(encoded)

# this model maps an input to its reconstruction
autoencoder = Model(input=input_img, output=decoded)

# this model maps an input to its encoded representation
encoder = Model(input=input_img, output=encoded)

# create a placeholder for an encoded (32-dimensional) input
encoded_input = Input(shape=(encoding_dim,))
# retrieve the last layer of the autoencoder model
decoder_layer = autoencoder.layers[-1]
# create the decoder model
decoder = Model(input=encoded_input, output=decoder_layer(encoded_input))

autoencoder.compile(optimizer='adadelta', loss='binary_crossentropy')
```

Using Theano backend.

Couldn't import dot_parser, loading of dot files will not be possible.

```
Using gpu device 0: Tesla K80 (CNMeM is disabled, CuDNN 5103)
/usr/prog/pythonML/3.0-goolf-1.5.14-NX-python-2.7.11/lib/python2.7
/site-packages/Theano-0.8.0-py2.7.egg/theano/sandbox/cuda/__init__
.py:600: UserWarning: Your CuDNN version is more recent than Theano. If you see problems, try updating Theano or downgrading CuDNN to version 4.
  warnings.warn(warn)
```

2. Load MNIST data

In [2]:

```
from keras.datasets import mnist
import numpy as np
(x_train, y_train), (x_test, y_test) = mnist.load_data()
```

3. Normalize images

In [3]:

```
x_train = x_train.astype('float32') / 255.  
x_test = x_test.astype('float32') / 255.  
x_train = x_train.reshape((len(x_train), np.prod(x_train.shape[1:])))  
x_test = x_test.reshape((len(x_test), np.prod(x_test.shape[1:])))  
print x_train.shape  
print x_test.shape
```

(60000, 784)

(10000, 784)

4. Train model

In [4]:

```
autoencoder.fit(x_train, x_train,  
                nb_epoch=50,  
                batch_size=256,  
                shuffle=True,  
                validation_data=(x_test, x_test))
```

WARNING (theano.gof.cmodule): WARNING: your Theano flags `gcc.cxxflags` specify an `-march=X` flags.

It is better to let Theano/g++ find it automatically, but we don't do it now

Train on 60000 samples, validate on 10000 samples

Epoch 1/50

60000/60000 [=====] - 1s - loss: 0.3750 -
val_loss: 0.2742

Epoch 2/50

60000/60000 [=====] - 1s - loss: 0.2683 -
val_loss: 0.2595

Epoch 3/50

60000/60000 [=====] - 1s - loss: 0.2492 -
val_loss: 0.2360

Epoch 4/50

60000/60000 [=====] - 1s - loss: 0.2270 -
val_loss: 0.2162

Epoch 5/50

60000/60000 [=====] - 1s - loss: 0.2103 -
val_loss: 0.2024

Epoch 6/50

60000/60000 [=====] - 1s - loss: 0.1987 -
val_loss: 0.1924

Epoch 7/50

60000/60000 [=====] - 1s - loss: 0.1900 -
val_loss: 0.1847

Epoch 8/50

60000/60000 [=====] - 1s - loss: 0.1829 -
val_loss: 0.1783

Epoch 9/50

60000/60000 [=====] - 1s - loss: 0.1768 -
val_loss: 0.1726

Epoch 10/50

```
60000/60000 [=====] - 1s - loss: 0.1714 -  
val_loss: 0.1674  
Epoch 11/50  
60000/60000 [=====] - 1s - loss: 0.1665 -  
val_loss: 0.1627  
Epoch 12/50  
60000/60000 [=====] - 1s - loss: 0.1620 -  
val_loss: 0.1585  
Epoch 13/50  
60000/60000 [=====] - 1s - loss: 0.1579 -  
val_loss: 0.1545  
Epoch 14/50  
60000/60000 [=====] - 1s - loss: 0.1541 -  
val_loss: 0.1510  
Epoch 15/50  
60000/60000 [=====] - 1s - loss: 0.1507 -  
val_loss: 0.1477  
Epoch 16/50  
60000/60000 [=====] - 1s - loss: 0.1476 -  
val_loss: 0.1447  
Epoch 17/50  
60000/60000 [=====] - 1s - loss: 0.1448 -  
val_loss: 0.1420  
Epoch 18/50  
60000/60000 [=====] - 1s - loss: 0.1421 -  
val_loss: 0.1394  
Epoch 19/50  
60000/60000 [=====] - 1s - loss: 0.1397 -  
val_loss: 0.1370  
Epoch 20/50  
60000/60000 [=====] - 1s - loss: 0.1374 -  
val_loss: 0.1348  
Epoch 21/50  
60000/60000 [=====] - 1s - loss: 0.1352 -  
val_loss: 0.1326  
Epoch 22/50  
60000/60000 [=====] - 1s - loss: 0.1332 -  
val_loss: 0.1306  
Epoch 23/50  
60000/60000 [=====] - 1s - loss: 0.1312 -  
val_loss: 0.1287  
Epoch 24/50  
60000/60000 [=====] - 1s - loss: 0.1293 -  
val_loss: 0.1268  
Epoch 25/50  
60000/60000 [=====] - 1s - loss: 0.1275 -  
val_loss: 0.1250  
Epoch 26/50  
60000/60000 [=====] - 1s - loss: 0.1258 -  
val_loss: 0.1234  
Epoch 27/50  
60000/60000 [=====] - 1s - loss: 0.1241 -  
val_loss: 0.1217  
Epoch 28/50  
60000/60000 [=====] - 1s - loss: 0.1225 -  
val_loss: 0.1202  
Epoch 29/50
```

```
60000/60000 [=====] - 1s - loss: 0.1210 -  
val_loss: 0.1187  
Epoch 30/50  
60000/60000 [=====] - 1s - loss: 0.1196 -  
val_loss: 0.1173  
Epoch 31/50  
60000/60000 [=====] - 1s - loss: 0.1182 -  
val_loss: 0.1160  
Epoch 32/50  
60000/60000 [=====] - 1s - loss: 0.1170 -  
val_loss: 0.1147  
Epoch 33/50  
60000/60000 [=====] - 1s - loss: 0.1158 -  
val_loss: 0.1135  
Epoch 34/50  
60000/60000 [=====] - 1s - loss: 0.1146 -  
val_loss: 0.1124  
Epoch 35/50  
60000/60000 [=====] - 1s - loss: 0.1136 -  
val_loss: 0.1114  
Epoch 36/50  
60000/60000 [=====] - 1s - loss: 0.1126 -  
val_loss: 0.1104  
Epoch 37/50  
60000/60000 [=====] - 1s - loss: 0.1117 -  
val_loss: 0.1095  
Epoch 38/50  
60000/60000 [=====] - 1s - loss: 0.1108 -  
val_loss: 0.1087  
Epoch 39/50  
60000/60000 [=====] - 1s - loss: 0.1100 -  
val_loss: 0.1079  
Epoch 40/50  
60000/60000 [=====] - 1s - loss: 0.1092 -  
val_loss: 0.1072  
Epoch 41/50  
60000/60000 [=====] - 1s - loss: 0.1085 -  
val_loss: 0.1065  
Epoch 42/50  
60000/60000 [=====] - 1s - loss: 0.1079 -  
val_loss: 0.1059  
Epoch 43/50  
60000/60000 [=====] - 1s - loss: 0.1072 -  
val_loss: 0.1052  
Epoch 44/50  
60000/60000 [=====] - 1s - loss: 0.1067 -  
val_loss: 0.1047  
Epoch 45/50  
60000/60000 [=====] - 1s - loss: 0.1061 -  
val_loss: 0.1041  
Epoch 46/50  
60000/60000 [=====] - 1s - loss: 0.1056 -  
val_loss: 0.1036  
Epoch 47/50  
60000/60000 [=====] - 1s - loss: 0.1051 -  
val_loss: 0.1032  
Epoch 48/50
```

```
60000/60000 [=====] - 1s - loss: 0.1047 -  
val_loss: 0.1027  
Epoch 49/50  
60000/60000 [=====] - 1s - loss: 0.1042 -  
val_loss: 0.1023  
Epoch 50/50  
60000/60000 [=====] - 1s - loss: 0.1038 -  
val_loss: 0.1019
```

Out[4]:

```
<keras.callbacks.History at 0x2aab14b0dc50>
```

5. Encode and decode images from test set

First row in the resulting figure are the input images Second row in the resulting figure are the reconstructed images. The 32D codes are in `encoded_imgs`

In [5]:

```
# encode and decode some digits  
# note that we take them from the *test* set  
encoded_imgs = encoder.predict(x_test)  
decoded_imgs = decoder.predict(encoded_imgs)  
## Here we can pass our own encodngs to the decoder to predict
```


In [6]:

```
%matplotlib inline
import matplotlib.pyplot as plt
#plt.rcParams['figure.figsize'] = (10, 10)
#plt.rcParams['image.interpolation'] = 'nearest'
plt.rcParams['image.cmap'] = 'gray'

n = 10 # how many digits we will display
plt.figure(figsize=(20, 4))
for i in range(n):
    # display original
    ax = plt.subplot(2, n, i + 1)

    plt.imshow(x_test[i].reshape(28, 28))
    plt.gray()
    ax.get_xaxis().set_visible(False)
    ax.get_yaxis().set_visible(False)

    # display reconstruction
    ax = plt.subplot(2, n, i + 1 + n)
    plt.imshow(decoded_imgs[i].reshape(28, 28))
    plt.gray()
    ax.get_xaxis().set_visible(False)
    ax.get_yaxis().set_visible(False)
plt.show()
```




6. Visualize encoded vectors

The encoder generates one 32D vector per image. We can visualize the collection of vectors with, e.g., PCA (Note: Deeper architectures would allow us to map directly the images to 2D vectors; <https://blog.keras.io/building-autoencoders-in-keras.html> (<https://blog.keras.io/building-autoencoders-in-keras.html>)). Below we visualize the vectors generated for the test dataset only (10k images). In the scatter plot, one data point corresponds to one image. The colors are the actual digits annotations (which were NOT used during training). Even with this shallow architecture we see that the digits are NOT randomly scattered, but rather the autoencoder managed to extract some structure within the data. The 32D vectors (or the PCA components, if you wish) can be used as feature vectors for, e.g., training other machine-learning methods.

In [34]:

```
from sklearn.decomposition import PCA
pca = PCA(n_components=2)
encoded_imgs = encoder.predict(x_test);
nImages = encoded_imgs.size / 32;
print nImages
data = encoded_imgs.reshape(nImages,32);
pca.fit(data);
scores = pca.transform(data);
plt.rcParams['image.cmap'] = 'jet'
plt.figure(figsize=(6, 6))
plt.scatter(scores[:, 0], scores[:, 1], c=y_test)
plt.colorbar()
plt.show()
```

10000



7. Synthesize some images with the decoder

We'll take a reference encoding vector from the test set and add a noise vector from a zero-mean Gaussian with standard deviation σ . By varying σ we can add a stronger distortion to the encoding, which leads to some interesting synthetically generated patterns. We could also pass directly the noise vector to the decoder, which yields even weirder images.

In [36]:

```
#sigma = 0.01;
#sigma = 0.1
#sigma = 1;
sigma = 3;
#sigma = 10;

n = 10 # how many digits we will display
plt.figure(figsize=(10, 2))
plt.rcParams['image.cmap'] = 'gray'
for i in range(n):
    # synthesize digit from a real image distorted with a random vector
    ax = plt.subplot(2, n, i + 1)
    noise = sigma * np.random.randn(32) # 32 is the encoding dimension
    mu = encoded_imgs[1]; # some reference encoding, to not be so far away in
the latent space
    z = mu + noise.reshape(-1,32);
    digitSynth = decoder.predict(z);
    plt.imshow(digitSynth.reshape(28, 28))
    plt.gray()
    ax.get_xaxis().set_visible(False)
    ax.get_yaxis().set_visible(False)
plt.show()
```



In []: